

In []:

```
#DL_P1 : Linear regression by using Deep Neural network:
# Implement Boston housing price prediction problem by Linear regression u
# Use Boston House price prediction dataset.
```

In [1]:

```
import pandas as pd
import numpy as np
from sklearn import metrics
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

In [2]:

```
import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
boston = load_boston()
```

In [3]:

```
data = pd.DataFrame(boston.data)
data.head()
```

Out[3]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [4]:

```
#Adding the feature names to the dataframe
data.columns = boston.feature_names
data.head()
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	

In [5]:

```
#Adding target variable to dataframe
data['PRICE'] = boston.target
```

In [6]:

```
#Check the shape of dataframe
data.shape
```

Out[6]: (506, 14)

```
In [7]: data.columns
```

```
Out[7]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
              'PTRATIO', 'B', 'LSTAT', 'PRICE'],  
             dtype='object')
```

```
In [8]: data.dtypes
```

```
Out[8]: CRIM      float64  
        ZN        float64  
        INDUS     float64  
        CHAS      float64  
        NOX       float64  
        RM        float64  
        AGE       float64  
        DIS       float64  
        RAD       float64  
        TAX       float64  
        PTRATIO   float64  
        B         float64  
        LSTAT     float64  
        PRICE     float64  
        dtype: object
```

```
In [9]: # Identifying the unique number of values in the dataset  
data.nunique()
```

```
Out[9]: CRIM      504  
        ZN        26  
        INDUS     76  
        CHAS       2  
        NOX       81  
        RM       446  
        AGE      356  
        DIS      412  
        RAD       9  
        TAX      66  
        PTRATIO   46  
        B       357  
        LSTAT    455  
        PRICE    229  
        dtype: int64
```

```
In [10]: # Check for missing values
data.isnull().sum()
```

```
Out[10]: CRIM      0
          ZN       0
          INDUS   0
          CHAS    0
          NOX     0
          RM      0
          AGE     0
          DIS     0
          RAD     0
          TAX     0
          PTRATIO 0
          B       0
          LSTAT   0
          PRICE   0
          dtype: int64
```

```
In [13]: # See rows with missing values
data[data.isnull().any(axis=1)]
```

```
Out[13]:  CRIM  ZN  INDUS  CHAS  NOX  RM  AGE  DIS  RAD  TAX  PTRATIO  B  LSTAT  PRICE
```

```
In [14]: # Viewing the data statistics
data.describe()
```

```
Out[14]:
```

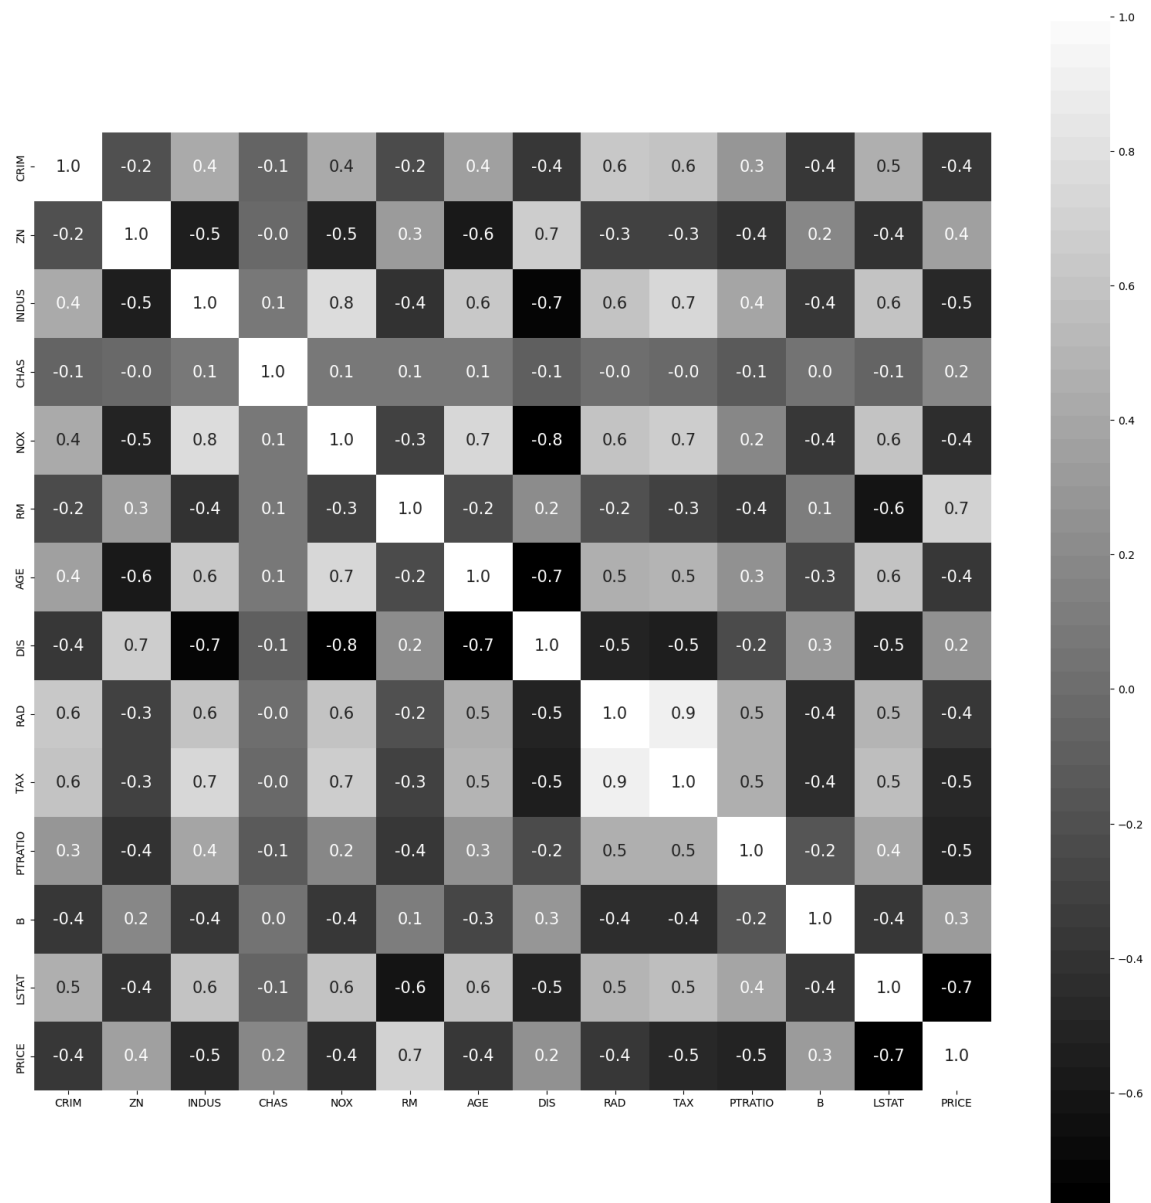
	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

```
In [15]: # Finding out the correlation between the features
corr = data.corr()
corr.shape
```

```
Out[15]: (14, 14)
```

```
In [16]: # Plotting the heatmap of correlation between features
plt.figure(figsize=(20,20))
sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True,
annot_kws={'size':15}, cmap='gray')
```

Out[16]: <AxesSubplot:>



```
In [17]: # Splitting target variable and independent variables
X = data.drop(['PRICE'], axis = 1)
y = data['PRICE']
```

```
In [18]: # Splitting to training and testing data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3,
random_state = 4)
```

```
In [19]: # Import library for Linear Regression
from sklearn.linear_model import LinearRegression
# Create a Linear regressor
lm = LinearRegression()
```

```
In [21]: # Train the model using the training sets
lm.fit(X_train, y_train)
```

Out[21]: LinearRegression()

```
In [22]: # Value of y intercept
lm.intercept_
```

Out[22]: 36.35704137659614

```
In [23]: #Converting the coefficient values to a dataframe
coefficients = pd.DataFrame([X_train.columns,lm.coef_]).T
coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coefficients'})
coefficients
```

Out[23]:

	Attribute	Coefficients
0	CRIM	-0.12257
1	ZN	0.055678
2	INDUS	-0.008834
3	CHAS	4.693448
4	NOX	-14.435783
5	RM	3.28008
6	AGE	-0.003448
7	DIS	-1.552144
8	RAD	0.32625
9	TAX	-0.014067
10	PTRATIO	-0.803275
11	B	0.009354
12	LSTAT	-0.523478

```
In [24]: # Model prediction on train data
y_pred = lm.predict(X_train)


# Model Evaluation
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train,
y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.089861094971133
MSE: 19.073688703469028
RMSE: 4.367343437774161
```


```
In [25]: # Visualizing the differences between actual prices and predicted values
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



```
In [26]: # Checking residuals  
plt.scatter(y_pred,y_train-y_pred)  
plt.title("Predicted vs residuals")  
plt.xlabel("Predicted")  
plt.ylabel("Residuals")  
plt.show()
```



```
In [27]: # Checking Normality of errors
sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```



```
In [28]: # Predicting Test data with the model
y_test_pred = lm.predict(X_test)
# Model Evaluation
acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test,
y_test_pred))*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7121818377409193
Adjusted R^2: 0.6850685326005711
MAE: 3.85900559237074
MSE: 30.053993307124163
RMSE: 5.482152251362978
```

```
In [ ]:
```